

# LAB 05: ABD QUORUMS

## ALGORITMOS E SISTEMAS DISTRIBUÍDOS (ASD)

**Alex Davidson**

October 2023

[a.davidson@fct.unl.pt](mailto:a.davidson@fct.unl.pt)

- ◇ Use your source code in Phase 1 of the project to implement a read/write register, on top of an ABD Quorum design.
- ◇ More details at <https://github.com/MEI-ASD-2023/lab05-abd-quorum>
- ◇ Will be useful for Phase 2 of the project

- ◇ Particular case of state-machine replication
- ◇ Only supports two functions: `read()`, `write(v)`
- ◇ `write(v)` returns ACK upon termination
- ◇ `read()` returns previously written value, or initial value  $v_0$

- ◇ Quorum-based system
- ◇ **Assumption (Liveness):**  $2f+1$  correct processes, for  $f$  possible failures
- ◇ **Assumption:** Asynchronous
- ◇ **Assumption:** Reliable (pp2p) links
- ◇ Generalises to multiple values

Identifiers for values, and parallel executions

State for process  $i$  ( $st_i$ ):

- ◇  $v_i$ : value, initially  $v_0$
- ◇  $t_i$ : pair  $(sn_i, id_i)$ , initially  $(0, 0)$
- ◇  $t_i > t_j$  iff  $(sn_i > sn_j) \vee ((sn_i = sn_j) \wedge (id_i > id_j))$

### Phase 1:

1. Send message `read-tag()` to all replicas
2. Wait for a quorum  $Q$  ( $|Q| > 2f + 1$ ) of replies
3. Let  $sn_{max} = \max(\{sn_i\}_{t_i \in Q})$

### Phase 2:

1. Send message `write( $t_v = (sn_{max} + 1, id_j), v$ )` to all
2. Wait for a quorum  $Q$  ( $|Q| > 2f + 1$ ) of ACK
3. Return ACK (write considered **terminated**)

**write( $v$ ) ISSUED BY PROCESS  $j$**


- ◇ Upon reception of `read-tag()`: return  $t_i$
- ◇ Upon reception of `write( $t_v, v$ )`:
  - ▷ if  $t_v > t_i$ :  $t_i = t_v$ ;  $v_i = v$
  - ▷ Return ACK
- ◇ Upon reception of `read()`: Return  $t_i, v_i$

1. Send `read()` to all replicas
2. Wait for a quorum  $Q$  of replies  $\{(t_i, v_i)\}_{i \in Q}$
3. Let  $(t_{max}, v_{max}) = \max(\{t_i, v_i\}_{i \in Q})$
4. Return  $v_{max}$

Question: Does this ensure **atomicity**?



Nope



This execution can occur because  $t_{max}$  may be derived from a non-completed write.

**Solution:** add a write-back( $p$ ) hase

1. **Phase 1:** Proceed as in naïve case
2. **Phase 2:** Send write( $t_{max}, v_{max}$ ) to all replicas
3. Wait for a majority quorum of ACK
4. Return  $v_{max}$

**Termination** ( $\leq f$  faults):

- ◇ Assuming well-formed messages
- ◇ Follows from properties of underlying links

- ◇ Serialisation points exist in **Phase 2** of `read()` and `write()`
- ◇ Writes necessarily happen **before** reads
- ◇ Reads return most recent value by a completed `write()`, otherwise `read()` **forces** `write()` completion.
- ◇ Subsequent `write()` ops will not be impacted by forced previous writes
- ◇ Serialisation points for `write()` are defined between start and end of operation

- ◇ “Sharing Memory Robustly in Message-Passing System” (JACM 1995)
- ◇ Edsger W. Dijkstra Prize in Distributed Computing: 2011
- ◇ Available [here](#)