

LAB 05: ABD QUORUMS

ALGORITMOS E SISTEMAS DISTRIBUÍDOS (ASD)

Alex Davidson

October 2023

a.davidson@fct.unl.pt

<https://classroom.github.com/a/JSdMJ6iF>

- ◇ Particular case of state-machine replication
- ◇ Only supports two functions: `read()`, `write(v)`
- ◇ `write(v)` returns ACK upon termination
- ◇ `read()` returns previously written value, or initial value v_0

- ◇ Quorum-based system
- ◇ Assumption (Liveness): $2f+1$ correct processes, for f possible failures
- ◇ Assumption: Asynchronous
- ◇ Assumption: Reliable (pp2p) links
- ◇ Generalises to multiple values

Identifiers for values, and parallel executions

State for process i (st_i):

- ◇ v_i : value, initially v_0
- ◇ t_i : pair (sn_i, id_i) , initially $(0, 0)$
- ◇ $t_i > t_j$ iff $(sn_i > sn_j) \vee ((sn_i = sn_j) \wedge (id_i > id_j))$

Phase 1:

1. Send message `read-tag()` to all replicas
2. Wait for a quorum Q ($|Q| > 2f + 1$) of replies
3. Let $sn_{max} = \max(\{sn_i\}_{t_i \in Q})$

Phase 2:

1. Send message `write($t_v = (sn_{max} + 1, id_j)$, v)` to all
2. Wait for a quorum Q ($|Q| > 2f + 1$) of ACK
3. Return ACK (write considered **terminated**)

write(v) ISSUED BY PROCESS j

- ◇ Upon reception of `read-tag()`: return t_i
- ◇ Upon reception of `write(t_v, v)`:
 - ▷ if $t_v > t_i$: $t_i = t_v$; $v_i = v$
 - ▷ Return ACK
- ◇ Upon reception of `read()`: Return t_i, v_i

1. Send `read()` to all replicas
2. Wait for a quorum Q of replies $\{(t_i, v_i)\}_{i \in Q}$
3. Let $(t_{max}, v_{max}) = \max(\{t_i, v_i\}_{i \in Q})$
4. Return v_{max}

Question: Does this ensure **atomicity**?

Nope



This execution can occur because t_{max} may be derived from a non-completed write.

Solution: add a write-back(p) hase

1. Phase 1: Proceed as in naïve case
2. Phase 2: Send write(t_{max}, v_{max}) to all replicas
3. Wait for a majority quorum of ACK
4. Return v_{max}

Termination ($\leq f$ faults):

- ◇ Assuming well-formed messages
- ◇ Follows from properties of underlying links

- ◇ Serialisation points exist in Phase 2 of `read()` and `write()`
- ◇ Writes necessarily happen **before** reads
- ◇ Reads return most recent value by a completed `write()`, otherwise `read()` **forces** `write()` completion.
- ◇ Subsequent `write()` ops will not be impacted by forced previous writes
- ◇ Serialisation points for `write()` are defined between start and end of operation

- ◇ “Sharing Memory Robustly in Message-Passing System” (JACM 1995)
- ◇ Edsger W. Dijkstra Prize in Distributed Computing: 2011
- ◇ Available [here](#)